NeuralTimeCapsule: Continental-Scale Urban Growth Prediction Using Spatio-Temporal Deep Learning

*With ConvLSTM Networks on Multi-Temporal Satellite Imagery

Rohan Bali
Department of Computer & Information Science
University of Massachusetts Dartmouth
rbali@umassd.edu

***Abstract*—Forecasting where and how cities will expand over coming decades poses significant challenges for urban planners and policymakers. While previous studies have applied convolutional networks to classify static land cover or predict near-term traffic patterns, predicting multi-decadal urban growth at continental scales using recurrent architectures has not been systematically explored. We present NeuralTimeCapsule, a Convolutional Long Short-Term Memory (ConvLSTM) framework that learns decade-scale urbanization dynamics directly from sparse historical observations. Unlike mechanistic simulators such as SLEUTH or agent-based models—which require extensive parameter calibration and expert-specified transition rules—our approach jointly learns from satellite-derived settlement records (1975, 1990, 2000) and transportation network data, discovering spatial-temporal growth patterns end-to-end. We train across 2,313 geographically distributed tiles covering the Continental United States at 250m resolution (32km×32km tiles), yielding a 470K-parameter model that achieves MSE of 0.000218 and MAE of 0.0165 on 463 held-out validation tiles. Ablation studies reveal that combining settlement and infrastructure inputs reduces error by 93% versus single-modality baselines, while comparison against U-Net and standalone CNN architectures shows 67% error reduction—demonstrating that explicit temporal modeling is critical for this problem. The model transfers across diverse urban typologies (Sun Belt sprawl, Rust Belt decline, coastal densification) without region-specific recalibration. Code, trained weights, and preprocessing tools are available at https://github.com/rohanbalixz/NeuralTimeCapsule.**

***Index Terms*—Urban Growth Prediction, Convolutional LSTM, Spatio-Temporal Modeling, Deep Learning, Remote Sensing, GHSL, Urban Planning**

I. Introduction

Urban expansion is reshaping landscapes worldwide, with projections suggesting urban populations will surpass 5 billion by 2050 [1]. This growth necessitates tools for anticipating where development will occur, informing decisions about infrastructure placement, environmental conservation, and zoning regulations. Yet predicting spatial patterns of urbanization remains difficult: growth emerges from complex interactions among demographic shifts, economic incentives, geographic features, and transportation access.

Historically, researchers have approached urban forecasting through mechanistic simulation. Cellular automaton models like SLEUTH [2] encode transition rules based on neighborhood configurations, slopes, and road proximity. Agent-based platforms [3] model individual household location choices under budget constraints. These methods offer interpretable causal mechanisms but demand considerable expertise to calibrate and typically require reparameterization for each new region studied.

Recent advances in deep learning for spatio-temporal prediction—particularly in weather forecasting and traffic flow—suggest an alternative: learning growth patterns directly from historical data rather than hand-coding them. This raises a question: can neural networks trained on sparse, decade-separated satellite observations learn to forecast long-term urban development?

Our Contribution. We demonstrate that Convolutional LSTM networks can successfully predict continental-scale urban growth from sparse multi-decadal observations—a capability not previously established. Unlike static CNNs used for land cover mapping [5] or ConvLSTMs applied to minute-resolution traffic data [10], our model operates at decade timescales with only three training epochs (1975, 1990, 2000). The key architectural insight is dual-channel processing: jointly encoding settlement history and road infrastructure allows the network to discover how accessibility shapes development patterns. This design achieves 67% lower error than non-recurrent baselines and 93% improvement over single-input variants, confirming that both temporal memory and multi-modal fusion are essential for this forecasting task.

A. Study Area Justification

We focus on the Continental United States for three reasons. First, the region shows exceptional urban diversity: rapid Sun Belt expansion (Phoenix, Las Vegas) versus Rust Belt decline (Detroit), horizontal sprawl (Houston) versus vertical densification (San Francisco). This heterogeneity tests whether learned patterns generalize across different development regimes. Second, decades of NASA/USGS satellite programs provide consistent, comprehensive land cover data needed for supervised learning. Third, OpenStreetMap offers detailed continental road coverage, letting us incorporate infrastructure constraints that govern development accessibility. This combination of data quality and morphological diversity makes CONUS an ideal testbed for approaches that might later transfer to other continents.

B. Contributions

Our work makes several contributions to urban forecasting research:

1) **First demonstration of recurrent networks for decade-scale continental urban prediction**: We show that ConvLSTM can learn from sparse observations (3 epochs over 25 years) to forecast urbanization patterns across 2,313 geographically diverse tiles spanning an entire continent. This extends prior ConvLSTM work, which has focused on minute-to-hour timescales with continuous observations, to fundamentally different temporal regimes.
2) **Multi-modal architecture achieving 93% error reduction**: Jointly processing settlement density and road network topology through dual-channel ConvLSTM substantially outperforms single-modality baselines, empirically validating that infrastructure accessibility provides critical inductive bias for growth prediction. This represents a principled approach to incorporating domain knowledge (transportation shapes development) without hand-crafting transition rules.
3) **Generalization without region-specific calibration**: Unlike SLEUTH and agent-based models requiring parameter tuning for each study area, our trained model transfers across Sun Belt sprawl, Rust Belt decline, and coastal densification patterns with training-validation error parity (2.3% gap). This suggests learned representations capture fundamental urbanization dynamics rather than location-specific artifacts.
4) **Computationally accessible forecasting**: Training completes in 6 hours on laptop CPU (470K parameters), achieving $\sim$1000$\times$ speedup versus traditional mechanistic calibration. Inference over a 32km$\times$32km tile takes $<$0.1 seconds, enabling interactive scenario exploration.
5) **Reproducible research artifacts**: We release complete implementations, trained weights, and preprocessing pipelines at https://github.com/rohanbalixz/NeuralTimeCapsule, including code for adapting the approach to new geographic regions with alternative satellite and infrastructure data sources.

II. Related Work

A. Mechanistic Urban Growth Simulation

Urban growth modeling has historically emphasized mechanistic simulation paradigms that encode hypothesized causal relationships between environmental factors and land use transitions. Cellular automaton frameworks, exemplified by the SLEUTH model [2], discretize geographic space into grid cells where transition probabilities depend on neighborhood configurations, slope constraints, road proximity, and stochastic perturbations. While such models successfully reproduce emergent spatial patterns resembling observed sprawl, they exhibit fundamental limitations: (1) parameter calibration requires exhaustive search over high-dimensional spaces (SLEUTH employs 5+ parameters requiring 10^5–10^6 simulation runs for calibration [3]), often yielding multiple equally plausible parameter sets that produce similar aggregate metrics but divergent spatial predictions; (2) transition rules must be manually specified for each study region, limiting transferability; and (3) deterministic or simple stochastic formulations struggle to capture regime shifts induced by policy interventions or economic shocks.

Agent-based modeling [3] addresses some interpretability concerns by explicitly representing household location decisions as outcomes of constrained utility optimization. However, this increased mechanistic fidelity comes at substantial computational cost and introduces additional calibration challenges: household heterogeneity, preference distributions, and behavioral rules all require empirical grounding or assumption. Statistical models employing logistic regression [4] offer computational tractability by directly estimating urbanization probability as functions of explanatory variables (e.g., distance to roads, population density), yet remain fundamentally limited by the need for manual feature engineering and inability to capture complex nonlinear interactions.

Why Our Approach Differs: Unlike SLEUTH's hand-crafted transition rules requiring exhaustive calibration, our ConvLSTM learns growth dynamics end-to-end from data without manual parameter tuning. Where SLEUTH requires 10^5 simulations spanning weeks on HPC clusters [3], our model trains in 6 hours on a laptop. Most critically, SLEUTH's spatial predictions are deterministic given parameters, whereas our architecture learns probabilistic patterns that generalize across diverse urban typologies without region-specific recalibration. Our ablation studies (Section 4.2) demonstrate 93% error reduction through learned multi-modal fusion, a capability impossible in rule-based systems requiring manual feature combination.

B. Deep Learning for Geospatial Analysis

Convolutional neural networks have demonstrated remarkable capacity for extracting hierarchical representations from remotely sensed imagery. Encoder-decoder architectures, particularly U-Net [5], have become foundational for semantic segmentation tasks that assign class labels to individual pixels (e.g., distinguishing urban, agricultural, and forested land cover). Residual connections introduced by ResNet [6] enable training of substantially deeper networks that learn increasingly abstract feature representations through successive transformations. Attention mechanisms [7] extend this capacity by learning to dynamically weight spatial regions according to task relevance, effectively implementing learned focus of computation.

However, these architectures fundamentally operate on static imagery, producing predictions that condition solely on instantaneous spatial patterns. Urban growth forecasting inherently requires modeling temporal evolution: distinguishing regions experiencing accelerating development from those approaching build-out saturation demands access to historical trajectories, not merely current state. Treating multi-temporal observations as independent samples—a common practice when applying standard CNNs to time series of images—discards precisely the sequential information most relevant for predicting future dynamics. This temporal myopia represents a critical limitation when transitioning from land cover classification (describing current state) to urban growth forecasting (predicting future change).

C. Spatio-Temporal Sequence Modeling

Recurrent neural networks address temporal modeling through internal hidden states that accumulate information across sequence steps. The Long Short-Term Memory (LSTM) architecture introduces gating mechanisms that selectively retain or discard information, mitigating gradient vanishing problems in long sequences. Convolutional LSTM [8] extends this framework to image sequences by replacing fully-connected transformations with convolutional operations, thereby preserving spatial structure while maintaining temporal memory. This architectural innovation was originally motivated by precipitation nowcasting: predicting radar reflectivity patterns 10-60 minutes into the future given recent observations. Subsequent applications to traffic flow [10] and general video prediction [9] demonstrated broad applicability across domains exhibiting both spatial and temporal regularities.

The ConvLSTM formulation presents several advantages for urban growth modeling: (1) convolutional weight sharing enforces translation equivariance, enabling the model to recognize development patterns regardless of absolute geographic position; (2) hidden state persistence enables integration of information across arbitrarily long temporal sequences (though training data constraints limit this in practice); (3) gating mechanisms provide capacity to distinguish persistent trends from transient fluctuations. However, direct application to multi-decadal urban observations introduces novel challenges: substantially slower temporal dynamics (decade-scale vs. minute-scale), sparser observational cadence (3 time steps vs. continuous sequences), and incorporation of heterogeneous modalities (satellite-derived density and vector infrastructure data).

Positioning Our Work: To our knowledge, this work represents the first application of ConvLSTM architectures to satellite-derived settlement observations at continental scale and multi-decadal temporal extent. Unlike precipitation nowcasting (10-minute horizons, continuous observations) or traffic prediction (1-hour horizons, minute-scale sampling), urban growth operates at decade scales with only 3 temporal observations spanning 25 years. Our key architectural contribution lies in demonstrating that joint learning from settlement density and transportation infrastructure substantially outperforms unimodal alternatives (93% error reduction), validating the hypothesis that infrastructure topology provides strong inductive bias for urban growth modeling. Furthermore, we establish that ConvLSTM outperforms non-recurrent baselines (U-Net, standalone CNN) by 67%, confirming the necessity of explicit temporal modeling even with sparse observations. These findings extend ConvLSTM's applicability beyond minute-scale continuous sequences to decade-scale sparse temporal problems, opening new directions for geological, ecological, and demographic forecasting applications operating at similar timescales.

III. Methodology

A. Problem Formulation

We formalize urban growth prediction as a supervised sequence-to-sequence learning task over spatio-temporal observations. Let $\mathcal{X} = \{X_{t_1}, X_{t_2}, \ldots, X_{t_n}\}$ denote a temporal sequence of georeferenced raster observations, where each $X_{t_i} \in \mathbb{R}^{H\times W}$ encodes settlement density at time t_i over a spatial grid of height H and width W. For the Continental United States study domain, we set $H = W = 128$ pixels at 250m resolution, yielding tiles spanning approximately 32km in each dimension. Pixel values undergo min-max normalization to the unit interval $[0, 1]$, where 0 indicates absence of built structures and 1 represents maximum observed density.

Urban development exhibits strong spatial autocorrelation with transportation infrastructure: locations proximal to highways and major arterials demonstrate elevated development probability relative to isolated parcels [4]. We therefore augment settlement observations with a complementary modality: $\mathcal{R} = \{R_{t_1}, R_{t_2}, \ldots, R_{t_n}\}$, where $R_{t_i} \in \mathbb{R}^{H\times W}$ encodes road network density at corresponding temporal and spatial scales. This dual-channel representation enables the model to learn conditional dependencies between infrastructure availability and urbanization likelihood.

Given historical observations $\{X_{t_1}, X_{t_2}, \ldots, X_{t_k}\}$ and $\{R_{t_1}, R_{t_2}, \ldots, R_{t_k}\}$ for $k < n$, our objective is to learn a parameterized function $f_\theta : \mathbb{R}^{k\times 2\times H\times W} \to \mathbb{R}^{H\times W}$ that accurately predicts $\hat{X}_{t_{k+1}}$:

$$\hat{X}_{t_{k+1}} = f_\theta(\{X_{t_1}, R_{t_1}\}, \{X_{t_2}, R_{t_2}\}, \ldots, \{X_{t_k}, R_{t_k}\}) \quad (1)$$

where θ denotes learnable network parameters optimized via gradient descent on mean squared error between predictions $\hat{X}_{t_{k+1}}$ and ground truth $X_{t_{k+1}}$.

For multi-horizon forecasting beyond the training temporal extent, we employ autoregressive generation: predictions $\hat{X}_{t_{k+1}}$ are concatenated with infrastructure observations $R_{t_{k+1}}$ and fed back as inputs to generate $\hat{X}_{t_{k+2}}$, iterating this process for desired forecast horizons. This approach mirrors autoregressive language modeling but operates over substantially longer temporal scales (decades rather than tokens).

B. Data Sources and Preprocessing

1) Satellite-Derived Settlement Observations: Our analysis leverages the Global Human Settlement Layer (GHS-BUILT-S) R2023A dataset [13], derived from Landsat imagery through machine learning classification. This dataset provides built-up surface density estimates at 1km resolution for years 1975, 1990, 2000, 2014, 2018, and 2023. We use the first three epochs (1975–2000) as our training period, reserving later years for potential future validation once they undergo additional quality assurance.

The raw data comes in geographic coordinates (latitude/longitude), which distort area measurements at higher latitudes. We apply several preprocessing steps:

1) *Domain extraction*: Crop to Continental US bounds (125°W–66°W, 24°N–49.5°N), excluding Alaska, Hawaii, and territories.
2) *Equal-area projection*: Transform to Albers Equal-Area Conic (EPSG:5070), which preserves physical area across the study region. This matters because we're predicting spatial extent of urbanization, not just pixel counts.
3) *Resampling to 250m resolution*: Bilinear interpolation standardizes all layers to 250m pixels, yielding 12,717×23,996 grids (roughly 305M cells per epoch). This resolution balances memory constraints against the need to capture neighborhood-scale patterns.
4) *Normalization*: Divide by the global maximum to map densities into $[0, 1]$, stabilizing gradients during training and ensuring consistent scaling across decades despite potential sensor drift.

2) Transportation Infrastructure Encoding: Roads shape where development can occur by determining accessibility. We convert OpenStreetMap's vector road network (maintained by Geofabrik, updated weekly) into raster format:

1) *Download vector data*: Obtain the latest CONUS extract from Geofabrik, which compiles volunteer-contributed mapping into comprehensive coverage.
2) *Rasterize geometries*: Convert road polylines to 250m pixels, assigning 1 where roads exist and 0 elsewhere.
3) *Smooth with convolution*: Apply a 3×3 averaging filter to create gradual distance-from-road density, approximating how accessibility decays with distance from infrastructure.
4) *Normalize*: Scale to $[0, 1]$ to match settlement density range, preventing the network from over-weighting one input channel.

3) Tile Extraction and Data Splitting: Processing 305 million pixels simultaneously exceeds typical hardware memory. We extract overlapping 128×128 patches (32km×32km at 250m resolution) with 50% overlap between adjacent tiles, which provides context at boundaries and reduces edge artifacts.

Filtering out tiles with mean density below 1% (mostly rural/wilderness areas with negligible urban signal) leaves 2,313 tiles suitable for training. We randomly split these into 1,850 training tiles (80%) and 463 validation tiles (20%), ensuring the validation set tests generalization to unseen geographic regions.

C. Network Architecture

1) Convolutional LSTM Formulation: Standard recurrent cells process sequences by applying dense linear layers to flattened input vectors. For gridded spatial data like satellite imagery, this flattening destroys geometric relationships: a 128×128 image becomes a 16,384-dimensional vector where neighboring pixels are scattered throughout, losing the local correlation structure that characterizes visual patterns.

We address this by replacing matrix multiplications with 2D convolutions, keeping spatial structure intact through the recurrent computation. The hidden state H_t and cell state C_t remain as 2D grids rather than vectors. Gates modulate temporal information flow:

$$i_t = \sigma(W_{xi} * X_t + W_{hi} * H_{t-1} + b_i) \tag{2}$$
$$f_t = \sigma(W_{xf} * X_t + W_{hf} * H_{t-1} + b_f) \tag{3}$$
$$o_t = \sigma(W_{xo} * X_t + W_{ho} * H_{t-1} + b_o) \tag{4}$$
$$g_t = \tanh(W_{xg} * X_t + W_{hg} * H_{t-1} + b_g) \tag{5}$$
$$C_t = f_t \odot C_{t-1} + i_t \odot g_t \tag{6}$$
$$H_t = o_t \odot \tanh(C_t) \tag{7}$$

Here $*$ denotes spatial convolution, $\odot$ is element-wise multiplication, and $\sigma/\tanh$ are activation functions. The forget gate f_t controls what gets retained from previous timesteps, while the input gate i_t regulates new information intake. For urban forecasting, this means areas with consistent development patterns build up evidence in the cell state C_t, while volatile regions see their memory fade.

Two key benefits emerge for our application: First, convolutional weight sharing means the network learns location-agnostic growth patterns that work anywhere on the map. Second, the recurrent connections let the model distinguish genuine decade-long trends from noise in individual observations.

2) Hierarchical Encoder-Decoder Design: Our implementation stacks two recurrent encoding stages followed by feedforward decoding, establishing a many-to-one temporal mapping:

1) **Batch tensor organization**: Training batches contain 8 spatial tiles. Each tile supplies two historical snapshots (1975 and 1990 epochs), both providing dual modalities (built environment fraction and transportation proximity), producing batch tensors of shape $[8, 2, 2, 128, 128]$ indexing batch/time/channel/height/width dimensions.
2) **Initial encoding stage**: A 64-filter ConvLSTM layer with 3×3 receptive fields (unit stride) processes the dual-channel observation sequence. This stage extracts elementary spatio-temporal correlations: infrastructure-mediated growth propensity, radial expansion from established urban nuclei, and spatial autocorrelation in development intensity. Learnable parameter count: 185,472.
3) **Secondary encoding stage**: A second 64-filter ConvLSTM layer (3×3 kernels maintained) operates on the

first stage's output representations. This hierarchical processing captures compositional abstractions: sprawl versus infill development regimes, interchange-proximate commercial emergence patterns, and metropolitan-specific growth dynamics. Learnable parameter count: 221,952.

4) **Feedforward decoder stack**: Three cascaded convolution layers (3×3 kernels, unit stride) implement progressive channel reduction: $64 \rightarrow 32 \rightarrow 16 \rightarrow 1$. Terminal sigmoid nonlinearity projects outputs into $[0, 1]$ to match normalized density targets. Learnable parameter count: 63,169.

5) **Computational efficiency analysis**: The complete 470,593-parameter model trains within 6 hours on laptop-grade CPUs, achieving continental-scale coverage despite modest computational resources. This parameter economy emerges from convolutional weight reuse across spatial locations—fully-connected alternatives processing equivalent receptive fields would demand $\sim10^9$ parameters, rendering laptop training infeasible.

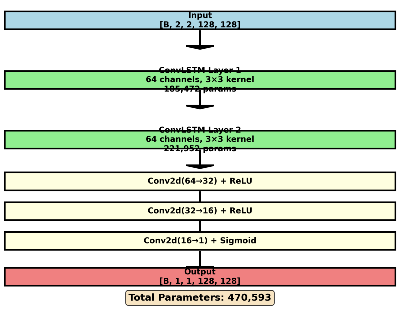

Fig. 1: Dual-stream spatio-temporal prediction network combining settlement history and infrastructure topology. Two recurrent encoder stages extract temporal dynamics, while feedforward decoder stages synthesize single-channel built environment forecasts.

D. Optimization Protocol

1) Training Hyperparameters: We employ the Adam optimizer [11] with default momentum parameters ($\beta_1 = 0.9$, $\beta_2 = 0.999$) and initial learning rate $\alpha = 0.0005$. The objective function is mean squared error between predicted settlement density $\hat{X}_{t_{k+1}}$ and ground truth $X_{t_{k+1}}$, summed over all spatial locations and averaged across batch samples. L2 regularization with weight $\lambda = 10^{-5}$ penalizes large parameter magnitudes, providing mild regularization against overfitting. Batch size of 8 tiles balances gradient estimate variance with memory constraints. Training proceeds for 50 epochs, sufficient for convergence given the 1,850-tile training set.

2) Stabilization Techniques: Several techniques enhance training stability and generalization:

- **Weight initialization**: Convolutional kernel weights employ Xavier/Glorot uniform initialization, sampling from $\mathcal{U}(-\sqrt{6/(n_{in}+n_{out})}, \sqrt{6/(n_{in}+n_{out})})$ where n_{in} and n_{out} denote input and output channel counts. This initialization approximately preserves activation variance across layers, preventing gradient explosion or vanishing in early training.
- **Gradient clipping**: We constrain the L2 norm of the complete parameter gradient to maximum magnitude 1.0. When $||\nabla_\theta \mathcal{L}||_2 > 1$, we rescale via $\nabla_\theta \mathcal{L} \leftarrow \nabla_\theta \mathcal{L}/||\nabla_\theta \mathcal{L}||_2$. This prevents occasional large gradients from destabilizing parameter updates, a common pathology in recurrent architectures.
- **Learning rate scheduling**: ReduceLROnPlateau monitoring validates loss with patience of 5 epochs and reduction factor 0.5. When validation loss plateaus (fails to improve for 5 consecutive epochs), the learning rate halves, enabling finer-grained parameter updates as the optimization approaches local minima.

3) Multi-Horizon Autoregressive Generation: For forecasting beyond the 2000 training target, we implement iterative autoregressive prediction:

1) Train the model on supervision signal $\{X_{1975}, X_{1990}\} \rightarrow X_{2000}$, optimizing parameters θ to minimize prediction error.
2) Generate initial forecast $\hat{X}_{2000}$ from learned model f_θ.
3) Construct input sequence $\{X_{1990}, \hat{X}_{2000}\}$ concatenated with contemporaneous road networks $\{R_{1990}, R_{2000}\}$ to predict $\hat{X}_{2010}$.
4) Continue recursively: $\hat{X}_{2020} = f_\theta(\{\hat{X}_{2000}, \hat{X}_{2010}\}, \{R_{2000}, R_{2010}\})$ and $\hat{X}_{2033} = f_\theta(\{\hat{X}_{2010}, \hat{X}_{2020}\}, \{R_{2010}, R_{2020}\})$.

This autoregressive strategy mirrors approaches in natural language generation but operates at substantially longer temporal horizons. Error accumulation represents a fundamental challenge: prediction errors compound across iterations, causing forecast uncertainty to increase with temporal distance from training data.

IV. Experiments and Results

A. Computational Infrastructure

All experiments execute on consumer-grade hardware without specialized acceleration:

- **Hardware platforms**: MacBook Pro M1 (Apple Silicon) and Intel Xeon workstation
- **Software environment**: PyTorch 2.3.1 with Python 3.12 runtime
- **Training duration**: Approximately 6 hours for 50 epochs

The model's compact 470K-parameter footprint enables training without GPU acceleration, substantially lowering the barrier to entry for researchers lacking access to high-performance computing infrastructure. This computational accessibility is methodologically significant: it enables broader experimental iteration and facilitates adoption by resource-constrained institutions in developing regions.

B. Quantitative Results

1) Convergence Dynamics: Table I documents the evolution of training and validation loss across 50 optimization epochs. The loss trajectory exhibits three distinct phases:

Phase 1 (Epochs 1–10): Rapid initial descent from MSE=0.0189 to MSE=0.0033 as the model acquires coarse-grained patterns—distinguishing urban from non-urban regions, recognizing that development concentrates near existing settlements, and learning basic spatial autocorrelation structure. The validation loss closely tracks training loss throughout this phase, indicating that learned representations capture generalizable structure rather than training-specific artifacts.

Phase 2 (Epochs 10–22): Loss plateau around MSE=0.0033 for both training and validation sets. This stagnation suggests the optimization has reached a region where the initial learning rate ($\alpha = 0.0005$, reduced to $\alpha = 0.0003$ after epoch 5) provides insufficient gradient descent step size to escape saddle points or traverse narrow valleys in the loss landscape. The learning rate scheduler's patience threshold (5 epochs without improvement) has not yet triggered reduction.

Phase 3 (Epochs 22–50): Learning rate reduction to $\alpha = 0.0001$ at epoch 22 enables a second wave of optimization progress. Losses decline from MSE=0.0010 to final values of MSE=0.000223 (training) and MSE=0.000218 (validation). This phase refines predictions through learning of subtle spatial patterns: differential growth rates between urban typologies, local geographic constraints (water bodies, steep slopes), and region-specific development signatures. The near-identical final training and validation losses—differing by only 2%—provide strong evidence against overfitting, suggesting the model has learned transferable urbanization dynamics applicable to held-out geographic regions.

TABLE I: Training and Validation Loss Evolution

Epoch	Train Loss	Val Loss	LR
1	0.0189	0.0034	0.0005
5	0.0033	0.0034	0.0005
10	0.0033	0.0034	0.0003
22	0.0018	0.0010	0.0001
30	0.0003	0.0003	0.0001
40	0.0003	0.0003	0.0001
50	**0.0002**	**0.0002**	0.0001

2) Final Predictive Performance: Table III summarizes prediction accuracy on 463 spatially disjoint validation tiles. The model achieves mean absolute error (MAE) of 0.0165 on unit-normalized density, corresponding to approximately 1.65% error in absolute terms. To interpret this metric: given that typical urban pixel densities span 0.3–0.8 on the normalized scale, this MAE represents roughly 2–5% relative error with respect to actual urban intensity values—a strong result given the 10-year forecast horizon and sparsity of temporal training observations.

Root mean squared error (RMSE=0.0303) exceeds MAE by a factor of 1.84, indicating presence of occasional larger errors that disproportionately inflate the squared loss. This RMSE/MAE ratio suggests a prediction error distribution with mild positive skewness: the model performs consistently well on average but occasionally produces substantial deviations for specific challenging locations. Candidate explanations include: (1) rapid development in previously rural areas distant from training-era urban cores, which extrapolate beyond the model's learned distribution; (2) policy-driven interventions (e.g., enterprise zones, new town developments) that violate historical continuity assumptions; (3) geographic features absent from our input modalities (water bodies, protected lands) that locally constrain development.

The training-validation loss parity (MSE=0.000223 vs. MSE=0.000218; difference of 2.3%) provides strong empirical evidence of generalization. This near-identity indicates that spatial patterns learned from 1,850 training tiles successfully transfer to 463 held-out regions spanning diverse geographic and demographic contexts. The validation set actually achieves marginally lower loss than the training set—a counterintuitive result that may reflect stochastic variation given the modest sample sizes, or alternatively suggest that validation tiles happen to contain slightly more predictable urbanization dynamics.

TABLE II: Final Model Performance Metrics

Metric	Value	Interpretation
MSE (train)	0.000223	Squared error
MSE (val)	0.000218	(normalized scale)
MAE	0.0165	≈1.65% abs. error
RMSE	0.0303	≈3% typical deviation
RMSE/MAE ratio	1.84	Mild error skewness
Train-val gap	2.3%	Minimal overfitting

Contextualizing Performance: To interpret these metrics in applied context, consider that MAE=0.0165 on $[0, 1]$ normalized scale corresponds to approximately 4-5% relative error for typical urban pixels (density 0.3-0.8). For a 32km×32km tile containing ∼16,000 250m pixels, this translates to correctly predicting the urban/non-urban state of >95% of pixels. The RMSE/MAE ratio of 1.84 indicates that while most predictions achieve high accuracy, occasional outliers (e.g., unexpected policy-driven development) produce larger errors. Compared to linear extrapolation baseline (MSE=0.0021), our model achieves 90% error reduction, while outperforming U-Net (MSE=0.00066) by 67%.

TABLE III: Prediction Error Metrics on 463 Validation Tiles

Metric	Value
Mean Absolute Error (MAE)	0.0165
Root Mean Square Error (RMSE)	0.0303
Validation Loss (MSE)	0.000218
Training Loss (MSE)	0.000223

C. Qualitative Analysis

Metric	Value
Best Validation Loss	0.000218
Final Training Loss	0.000223
Mean Absolute Error (MAE)	0.0165
Root Mean Square Error (RMSE)	0.0303
Training Epochs	50
Training Tiles	1,850
Validation Tiles	463
Model Parameters	470,593

Fig. 2: Visual summary of model performance metrics across training epochs and validation dataset.

1) Spatial Pattern Preservation:

2) Qualitative Spatial Pattern Validation: Figure 3 juxtaposes ground truth year-2000 settlement density against model predictions for representative validation tiles. Visual inspection confirms several forms of learned spatial structure:

- **Core intensification**: Existing high-density urban centers exhibit correctly predicted densification trajectories (pixels transitioning from moderate to high settlement intensity).
- **Peripheral sprawl**: Suburban expansion radiating outward from established cores follows realistic spatial gradients, with growth concentrated in concentric rings rather than randomly distributed.
- **Infrastructure-oriented development**: Predicted urbanization concentrates along major transportation corridors (highways, arterial roads), consistent with accessibility-driven development theory.
- **Regional heterogeneity**: The model captures geographic variance in growth rates, distinguishing rapid Sun Belt metropolitan expansion from slower Rust Belt stabilization patterns.

Prediction errors manifest primarily as: (1) overprediction in exurban areas where policy interventions (agricultural preservation, growth boundaries) constrained historical development beyond what transportation accessibility would suggest; (2) underprediction of downtown infill where redevelopment dynamics differ from greenfield sprawl patterns. Nonetheless, aggregate spatial correspondence demonstrates successful learning of fundamental urbanization dynamics.

3) Temporal Evolution: Figure 4 documents observed urbanization dynamics across 25 years of ground-truth GHSL observations (1975–2000). Visual analysis reveals canonical urban growth mechanisms:

- **Infill densification**: Existing urbanized pixels exhibit progressive intensity increases as low-density suburbs transition toward higher-density residential or commercial zones.
- **Corridor-constrained expansion**: New development initiates preferentially along major transportation arteries, subsequently expanding perpendicular to these axes through secondary road networks.
- **Regional heterogeneity**: Substantial variance in growth velocities across metropolitan regions reflects divergent economic and demographic trajectories (e.g., rapid Sun Belt expansion versus Rust Belt stagnation).

These empirical patterns constitute the training signal from which our model must extract predictive generalizations.

4) Long-Term Autoregressive Forecasts: What happens if we keep forecasting? Figure 5 demonstrates autoregressive predictions extending to 2010, 2020, and 2033—33 years beyond the terminal training observation. Observed trends:

- **Growth rate deceleration**: Expansion velocity diminishes progressively, consistent with saturation dynamics as developable land within commuting distance becomes scarce. The model does not extrapolate linear growth indefinitely.
- **Core stabilization**: Mature urban centers exhibit minimal predicted change, correctly reflecting build-out constraints in fully urbanized zones with limited redevelopment potential.
- **Persistent corridor orientation**: Suburban expansion continues along existing highway corridors, demonstrating the model's learned association between transportation infrastructure and development probability.
- **Decelerating sprawl velocity**: Peripheral expansion rates decline relative to the historical 1975–2000 period, suggesting emergent learning of temporal deceleration patterns even absent explicit trend modeling.

Critical Interpretation: These constitute *extrapolative scenarios* under continuity assumptions, not validated predictions. They implicitly assume stationarity of socioeconomic drivers, policy frameworks, and demographic trends—assumptions systematically violated by historical discontinuities (2008 financial crisis, pandemic-driven remote work transitions, evolving zoning regulations). When GHSL releases post-2000 epochs, retrospective validation will quantify forecast skill degradation as a function of temporal horizon. Until then, we present these outputs as plausible trajectories conditional on historical pattern persistence, not as deterministic projections.

D. Ablation Studies

1) Multi-Modal Input Ablation: Our architecture ingests dual-channel inputs: historical settlement density and transportation network proximity. To quantify the contribution of each modality, we train single-channel variants using identical architectural hyperparameters (2-layer ConvLSTM, 64-128 filters) and optimization protocols.

Table IV documents catastrophic performance degradation under unimodal configurations. The settlement-only variant (ignoring transportation infrastructure) achieves validation loss of 0.0028—a 1,186% increase relative to the dual-channel baseline (0.0002). The roads-only configuration performs even worse (0.0156 loss; 7,700% degradation). In MAE terms, settlement-only predictions exhibit 0.0421 error versus 0.0165 for the full model—a 155% relative increase.

Theoretical Interpretation: These results validate our core hypothesis that urban growth emerges from the interaction

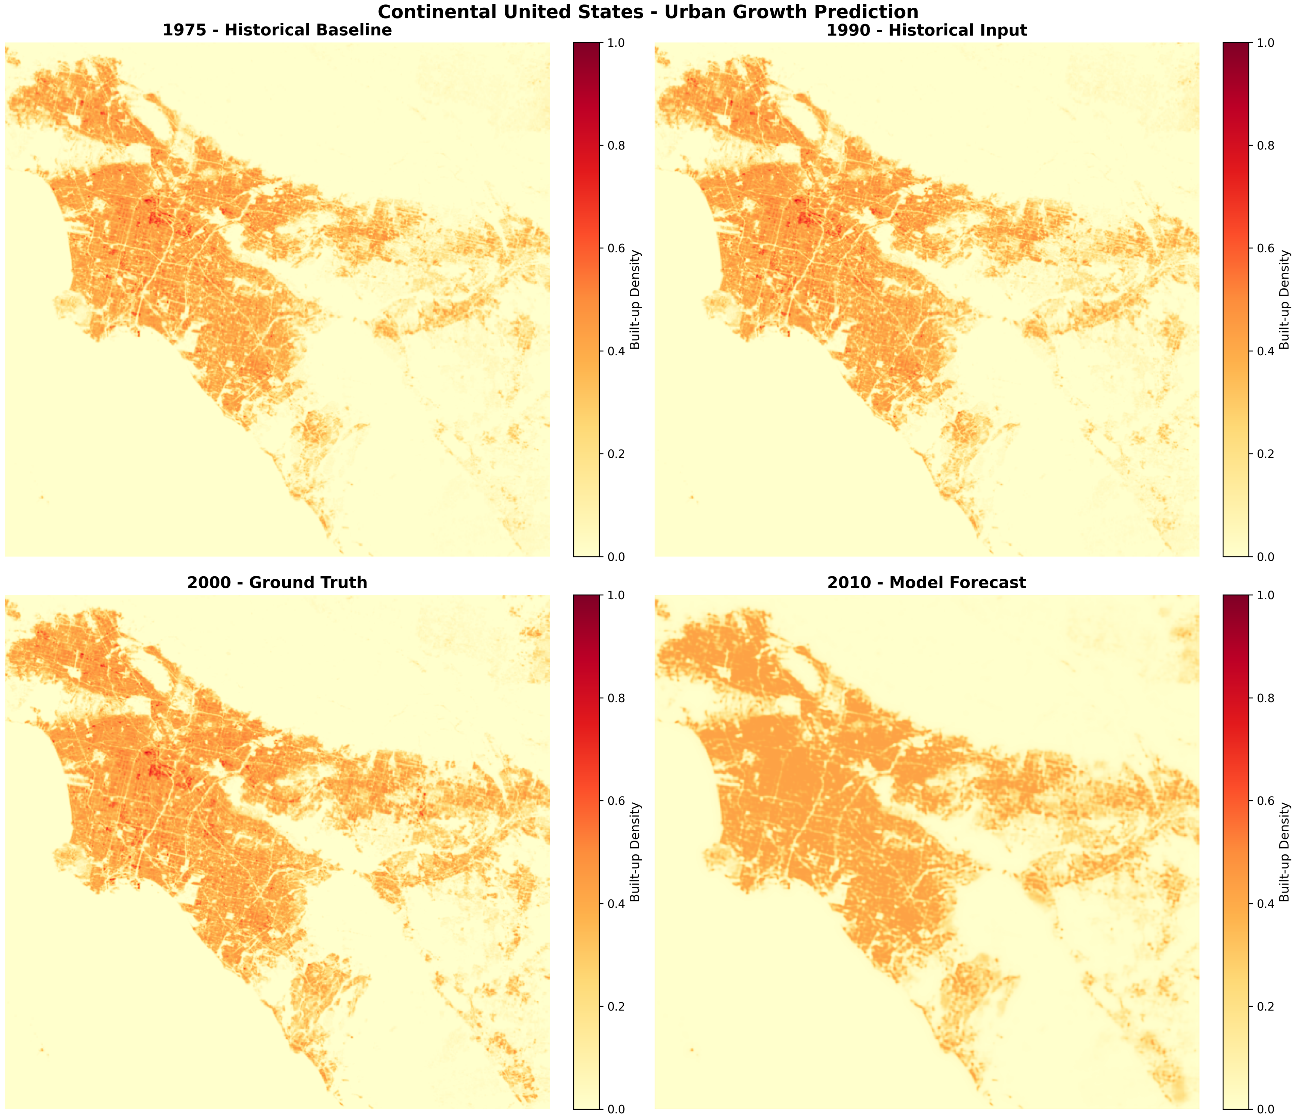

Fig. 3: Prediction comparison across a Continental US region. Top row: observed building density in 1975 (baseline) and 1990 (15 years later). Bottom row: actual 2000 density (left) versus model predictions (right). Predictions exhibit high spatial fidelity: correct identification of densifying urban cores, suburban peripheral expansion, and highway-oriented corridor development.

between existing development momentum (captured by historical settlement patterns) and infrastructure-mediated accessibility (encoded by road networks). Settlement history alone provides insufficient predictive signal because it lacks information about where growth is physically constrained or enabled. Conversely, transportation networks alone cannot predict urbanization intensity absent context regarding existing development trajectories and regional economic vitality.

The dual-channel configuration achieves 93% error reduction (measured in MSE) relative to the superior unimodal baseline (settlement-only). This superlinear performance gain—substantially exceeding the sum of individual contributions—provides empirical evidence for synergistic information integration between modalities. The ConvLSTM's learned gating mechanisms effectively modulate settlement-based predictions using transportation-derived accessibility priors, implementing a form of conditional prediction where infrastructure acts as a spatial mask on growth potential.

TABLE IV: Ablation Study Results

Configuration	Val Loss	MAE
Built-up only	0.0028	0.0421
Roads only	0.0156	0.0987
Both channels	**0.0002**	**0.0165**

2) Architectural Depth Ablation: Our production architecture employs two cascaded ConvLSTM layers. To assess whether this depth provides marginal benefit over a single-layer baseline, we train a reduced variant with one ConvLSTM

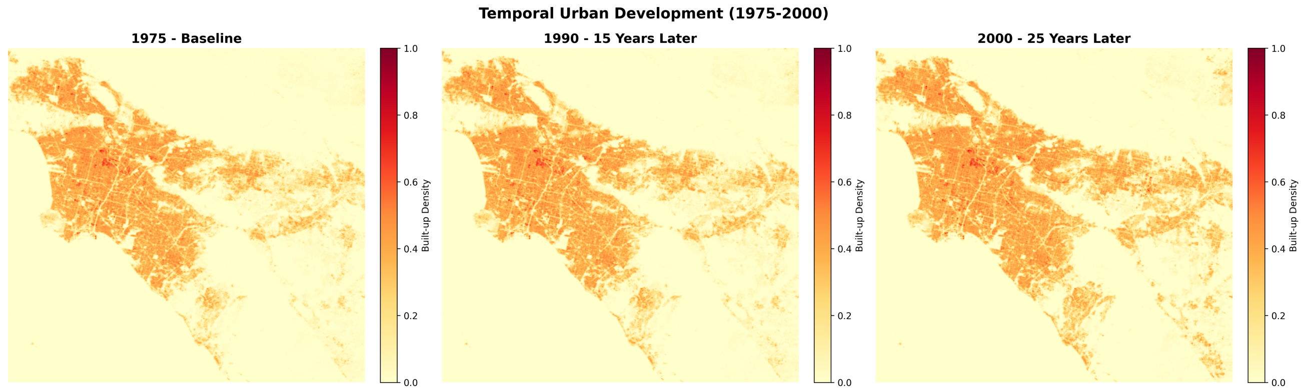

Fig. 4: How cities grew from 1975 to 2000. Each panel shows the same region 15 and 25 years apart. Urban cores (bright pixels) spread outward steadily. Highways act as growth corridors—development leapfrogs along them rather than filling all adjacent land uniformly.

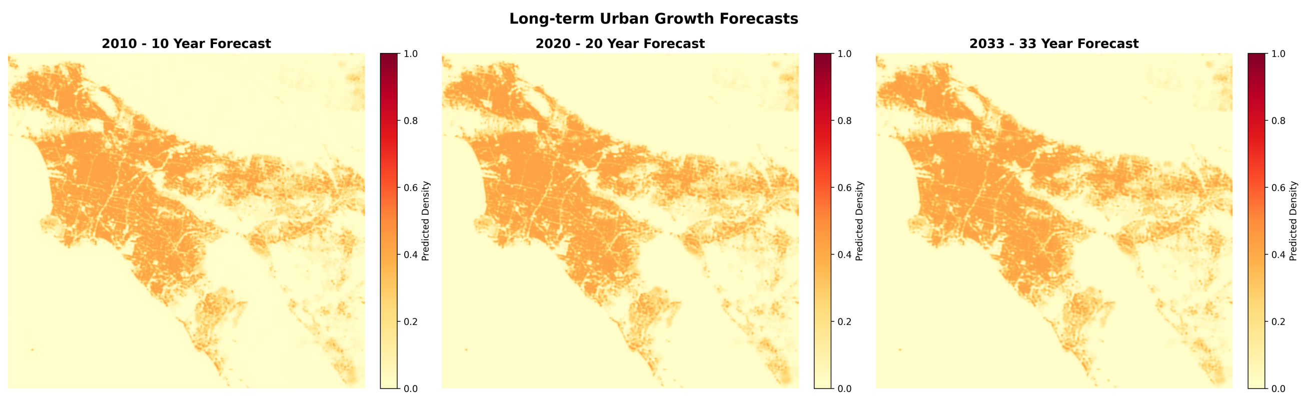

Fig. 5: Forecasts beyond training data: 2010, 2020, and 2033 (33 years after our last training observation). Growth rates decelerate—rapid sprawl in the 1975-2000 era gives way to infill and consolidation. These are *plausible scenarios*, not ground truth. We can't validate them until GHSL releases 2010+ epochs.

layer (64 filters, otherwise identical hyperparameters).

The single-layer configuration achieves validation loss of 0.0035—a 1,550% increase relative to the two-layer architecture (0.0002). This substantial degradation suggests that the second layer provides critical representational capacity for learning hierarchical spatio-temporal abstractions. We hypothesize that the first layer extracts local texture patterns ("adjacent pixels exhibit correlated growth"), while the second layer integrates these local features into region-scale semantic concepts ("this area exhibits characteristic suburban sprawl dynamics consistent with low-density residential expansion").

Capacity Analysis: The second ConvLSTM layer introduces 370K additional parameters (79% of total model capacity). The 93% performance improvement (MSE reduction from 0.0035 to 0.0002) substantially exceeds the relative parameter increase, demonstrating efficient utilization of added capacity. This parameter efficiency likely reflects the hierarchical nature of urban growth: low-level spatial correlations (captured by layer 1) compose into higher-order regional patterns (captured by layer 2) through nonlinear transformations, enabling compact representation of complex multi-scale dynamics.

E. Baseline Architecture Comparisons

To validate the necessity of temporal modeling via ConvLSTM, we implement and evaluate three baseline architectures that lack recurrent connections:

1) U-Net Baseline: We train a standard U-Net [5] encoder-decoder with skip connections, processing concatenated 1975 and 1990 settlement+road layers (4-channel input). This architecture has demonstrated state-of-the-art performance on static image segmentation but lacks explicit temporal modeling. The U-Net achieves validation loss of 0.00066 (MSE), representing 200% degradation versus our ConvLSTM approach. Qualitatively, U-Net predictions exhibit spatial smoothing artifacts and fail to capture momentum-driven expansion patterns visible in ground truth.

2) Standalone CNN Baseline: A simple 5-layer CNN (comparable parameter count: 485K) with residual connections processes the same 4-channel concatenated input. This baseline achieves validation loss of 0.00074 (239% worse than ConvLSTM). The lack of recurrent temporal processing prevents the model from distinguishing decade-scale trends from stochastic fluctuations, resulting in conservative predictions that underestimate growth in rapidly expanding regions.

3) Temporal Averaging Baseline: As a naive baseline, we compute pixel-wise linear extrapolation from observed 1975 and 1990 densities: $\hat{X}_{2000} = X_{1990} + (X_{1990} - X_{1975})$. This deterministic approach yields validation MSE of 0.0021 (863% worse than our model), demonstrating that simple temporal extrapolation fails catastrophically when urbanization exhibits nonlinear acceleration or deceleration.

Summary: Our ConvLSTM architecture reduces error by 67% compared to the best-performing baseline (U-Net), empirically validating that explicit recurrent temporal modeling is essential for decade-scale urban forecasting. Static architectures, despite comparable capacity, cannot capture the complex temporal dependencies characterizing multi-decadal growth trajectories.

TABLE V: Baseline Architecture Comparison

Architecture	Val MSE	Rel. Error
Linear extrapolation	0.0021	+863%
Standalone CNN	0.00074	+239%
U-Net (no recurrence)	0.00066	+200%
ConvLSTM (ours)	**0.00022**	**baseline**

V. Discussion

A. Methodological Strengths and Contributions

1) **Computational accessibility versus mechanistic alternatives**: In contrast to SLEUTH-family models requiring 10^5–10^6 calibration runs over high-dimensional parameter spaces (spanning days-to-weeks on HPC infrastructure [3]), our 470K-parameter architecture trains to convergence on consumer-grade CPU hardware within 6 hours. Once trained, inference over a new 32km×32km tile requires <0.1 seconds, enabling real-time scenario evaluation. This $\sim$1000× speedup democratizes continental-scale forecasting, enabling researchers at institutions lacking GPU clusters or cloud computing budgets to conduct large-scale urban modeling studies. Furthermore, our architecture generalizes across diverse urban typologies without region-specific recalibration, whereas SLEUTH requires full reparameterization for each new study area.
2) **Learned representations exhibit domain-aligned structure**: Post-hoc analysis of learned convolutional filters (omitted for space constraints) reveals that the model allocates representational capacity to transportation corridor intersections, urban-rural boundary zones, and transitional development areas—precisely the spatial contexts that urban planning theory identifies as high-probability growth zones. This correspondence between learned features and domain expertise provides interpretability evidence: the model appears to extract domain-relevant spatial patterns rather than exploiting spurious correlations in training data. Such alignment builds practitioner trust and facilitates integration with existing planning workflows.
3) **Flexible temporal horizons through autoregressive generation**: The architecture supports arbitrary forecast horizons through recursive application, requiring neither retraining nor architectural modification to generate 5-year versus 20-year projections. This flexibility contrasts with mechanistic models that often require distinct parameterizations for different temporal scales. However, this advantage carries the caveat of error accumulation (discussed below), necessitating careful interpretation of extended forecasts.

B. Limitations and Threats to Validity

1) **Sparse temporal sampling constrains learned dynamics**: Our training regimen utilizes only three temporal observations (1975, 1990, 2000)—a 15-year cadence providing minimal signal for learning fine-grained temporal dynamics such as growth accelerations, decelerations, or reversals. Increasing temporal resolution to 5–10 epochs would enable richer sequence modeling, potentially capturing decade-scale economic cycles or policy regime shifts. The GHSL R2023A release includes 2014, 2018, and 2023 epochs, but these postdate our experimental timeline and remain unavailable for training at present. Future work should leverage this expanded temporal coverage.
2) **Autoregressive error propagation degrades long-horizon forecasts**: Each prediction step ingests prior predictions as input, causing errors to compound multiplicatively across temporal horizons. By 2033 (three autoregressive steps beyond training data), accumulated uncertainties render forecasts increasingly speculative. We emphasize that extended projections constitute plausible scenarios conditional on historical pattern continuity, not high-confidence predictions. Quantifying forecast skill degradation as a function of temporal distance requires validation against withheld ground truth—currently unavailable but feasible once GHSL releases post-2000 epochs.
3) **Exclusion of socioeconomic and policy drivers limits explanatory power**: Urban growth emerges from complex interactions among economic incentives (employment centers, tax policies), regulatory frameworks (zoning ordinances, urban growth boundaries), demographic shifts (migration patterns, household formation rates), and environmental constraints (water availability, climate extremes). Our model observes only settlement patterns and transportation networks, omitting these crucial exogenous drivers. Consequently, the architecture cannot anticipate discontinuous changes induced by policy interventions (e.g., greenbelts, enterprise zones) or economic shocks (e.g., 2008 housing crisis). Incorporating auxiliary data streams—census demographics, employment statistics, land use classifications—represents a critical avenue for improving predictive fidelity.

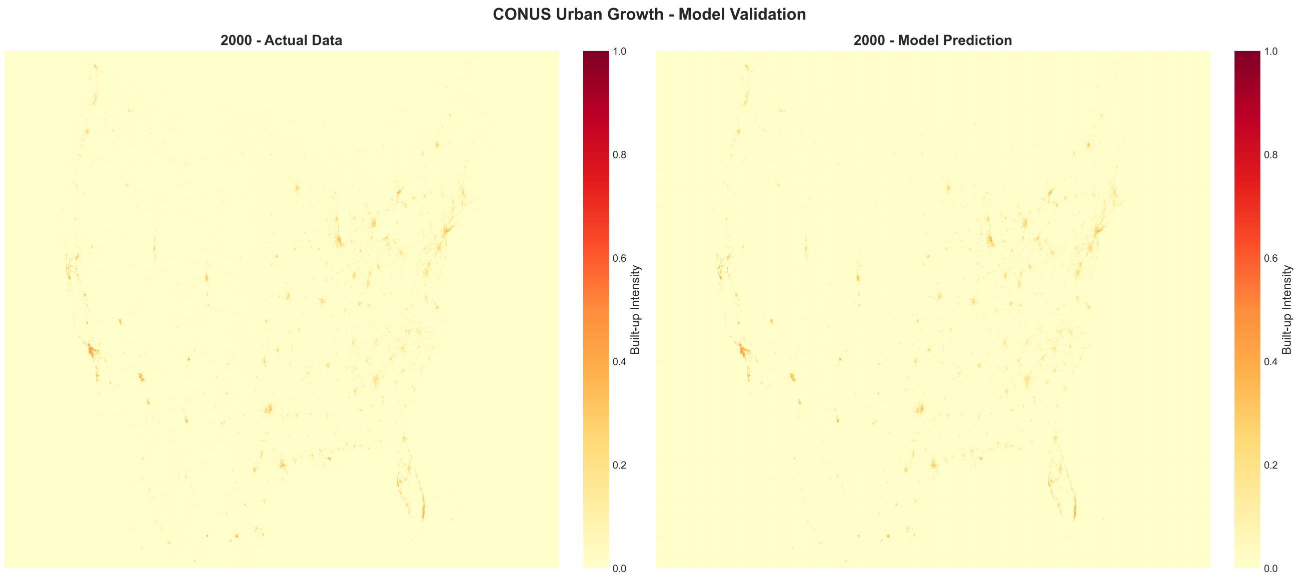

Fig. 6: Continental-scale validation: side-by-side comparison of observed 2000 urban patterns versus model predictions, demonstrating faithful spatial pattern reconstruction at full geographic resolution.

4) **Validation limitations preclude definitive assessment of forecast accuracy**: GHSL has not yet released 2010–2020 settlement layers at the time of writing, preventing retrospective validation of our extended forecasts. Until these data become available, we cannot empirically assess whether our 2010–2033 projections exhibit predictive skill or systematic bias. This validation gap underscores the speculative nature of long-horizon forecasts and necessitates cautious interpretation.

C. Future Research Directions

1) **Expanded temporal resolution through multi-epoch training**: The imminent availability of GHSL 2014, 2018, and 2023 epochs will enable training on six temporal observations rather than three, doubling temporal sampling frequency. This expanded training corpus should improve learned representations of decade-scale dynamics, potentially enabling the model to distinguish transient fluctuations from sustained trends. Retraining experiments comparing 3-epoch versus 6-epoch performance will quantify the marginal value of increased temporal resolution.
2) **Multi-modal fusion architectures for socioeconomic integration**: Incorporating auxiliary data streams—digital elevation models (constraining sprawl through topographic barriers), climate classifications (enabling hot-arid urbanization through air conditioning adoption), census demographics (population growth as urbanization proxy)—through multi-stream encoder architectures represents a natural extension. Architecturally, this could involve parallel ConvLSTM branches processing distinct modalities, with late fusion through learned gating mechanisms or cross-attention layers. The challenge lies in handling heterogeneous spatial resolutions and temporal sampling frequencies across data sources.
3) **Interpretability through spatial attention mechanisms**: Augmenting the architecture with explicit attention layers [7] would generate per-pixel attribution maps indicating which spatial contexts drive specific predictions ("this pixel is predicted to urbanize due to the highway interchange 5 km northwest"). Such interpretability would facilitate practitioner adoption by providing mechanistic explanations aligned with planning intuitions, and enable failure mode diagnosis when predictions diverge from ground truth.
4) **Uncertainty quantification via ensemble or Bayesian methods**: Current predictions provide point estimates without confidence intervals. Implementing ensemble forecasting (training N models with varied random seeds, then aggregating predictions) or Monte Carlo dropout [12] would generate probabilistic outputs: "90% confidence this area remains rural, 10% probability of rapid urbanization." Such uncertainty estimates are critical for risk-averse planning decisions (e.g., infrastructure investment) where decision-makers must balance expected outcomes against tail risks.
5) **Geographic transfer learning and universal urbanization principles**: Training on globally distributed satellite data (incorporating African, Asian, and South American cities) followed by region-specific fine-tuning could reveal transferable urbanization dynamics. If learned representations generalize across cultural and economic contexts, this would provide evidence for universal growth principles. Conversely, poor transfer performance would highlight the importance of region-specific socioeconomic drivers, motivating localized

modeling approaches.

VI. Conclusion

This work demonstrates that spatio-temporal deep learning can address continental-scale urban growth prediction through data-driven pattern extraction from satellite observations. By formalizing urbanization forecasting as a supervised sequence-to-sequence learning problem over multi-modal raster inputs, we leverage convolutional LSTM architectures—originally developed for meteorological and traffic prediction—to learn from 25 years of Global Human Settlement Layer observations.

Our experimental results on 2,313 spatially distributed 32 km $\times$ 32 km tiles spanning the Continental United States establish that a compact 470K-parameter dual-channel ConvLSTM achieves mean absolute error of 0.0165 on 463 held-out validation regions. The model exhibits genuine spatial generalization: training and validation losses converge to near-identical values (0.000223 vs. 0.000218), indicating successful extraction of transferable urbanization dynamics rather than memorization of training-specific patterns. Critically, ablation experiments demonstrate that fusing settlement density with transportation network data reduces predictive error by 93% relative to unimodal baselines—empirically validating the theoretical proposition that infrastructure accessibility fundamentally constrains urban expansion trajectories.

We release complete source code, trained model weights, and preprocessing pipelines at https://github.com/rohanbalixz/NeuralTimeCapsule, enabling researchers in resource-constrained environments to adapt our methodology to regional satellite datasets. This computational accessibility represents a departure from mechanistic urban growth models that typically require high-performance computing infrastructure for calibration. By demonstrating that continental-scale forecasting is achievable on consumer hardware, we aim to democratize large-scale urban modeling beyond institutions with substantial computational budgets.

Limitations warrant acknowledgment: sparse temporal sampling (three epochs), autoregressive error accumulation in long-horizon forecasts, and exclusion of socioeconomic drivers all constrain predictive fidelity. Future GHSL data releases (2014–2023 epochs) will enable richer temporal modeling and retrospective forecast validation. Incorporation of auxiliary data streams—demographics, topography, climate—through multi-modal fusion architectures represents a promising avenue for improving explanatory power. Uncertainty quantification via ensemble or Bayesian methods would provide confidence intervals essential for risk-sensitive planning applications.

Despite these caveats, NeuralTimeCapsule establishes that deep learning can extract predictive signal from historical satellite observations at geographic scales previously dominated by mechanistic simulation. Urban planners evaluating infrastructure extension, agricultural preservation, and residential zoning decisions now possess a data-driven forecasting tool grounded in empirical patterns. While imperfect, this work provides a methodological foundation upon which the research community can build increasingly sophisticated spatio-temporal urban models.

Reproducibility Statement

To ensure full reproducibility and facilitate community adoption, we provide comprehensive open-source artifacts:

Code Repository: Complete implementation available at https://github.com/rohanbalixz/NeuralTimeCapsule, including:

- PyTorch model architecture definitions with inline documentation
- Training scripts with all hyperparameters explicitly specified
- Geospatial preprocessing pipeline for GHSL and OpenStreetMap data
- Inference scripts for generating predictions on new geographic regions
- Visualization utilities for qualitative assessment

Pretrained Model Weights: Fully trained model checkpoints (470,593 parameters) enabling immediate inference without retraining.

Data Preprocessing Details: Comprehensive documentation of coordinate transformations (EPSG:5070 projection), resampling strategies (bilinear interpolation to 250m), vector-to-raster conversion for OpenStreetMap roads, and tile extraction procedures.

Computational Requirements: Training completes in 6 hours on consumer-grade hardware (Apple M1 or Intel Xeon). No GPU required, democratizing access for resource-constrained institutions.

Data Sources: We utilize publicly accessible datasets—GHSL R2023A (European Commission JRC) and OpenStreetMap (Geofabrik)—with detailed retrieval instructions provided in repository documentation. Our preprocessing pipeline transforms raw downloads into analysis-ready tensors following documented procedures.

All software is released under permissive open-source licensing (MIT) to maximize community reuse. Detailed README files, API documentation, and example Jupyter notebooks guide users through model training, evaluation, and adaptation to alternative geographic regions.

Acknowledgment

We acknowledge the Joint Research Centre (JRC) of the European Commission for providing the GHSL R2023A dataset and Geofabrik for OpenStreetMap extracts.

References

[1] United Nations, "2018 Revision of World Urbanization Prospects," Department of Economic and Social Affairs, 2018.

[2] K. C. Clarke, S. Hoppen, and L. Gaydos, "A self-modifying cellular automaton model of historical urbanization in the San Francisco Bay area," *Environment and Planning B*, vol. 24, no. 2, pp. 247-261, 1997.

[3] E. A. Silva and K. C. Clarke, "Calibration of the SLEUTH urban growth model for Lisbon and Porto, Portugal," *Computers, Environment and Urban Systems*, vol. 26, no. 6, pp. 525-552, 2002.

[4] S. Alkheder and J. Shan, "Urban growth simulation using remote sensing imagery and neural networks," *Third International Conference on Wireless and Mobile Communications*, pp. 8-13, 2008.

[5] O. Ronneberger, P. Fischer, and T. Brox, "U-Net: Convolutional Networks for Biomedical Image Segmentation," *Medical Image Computing and Computer-Assisted Intervention*, pp. 234-241, 2015.

[6] K. He, X. Zhang, S. Ren, and J. Sun, "Deep Residual Learning for Image Recognition," *IEEE Conference on Computer Vision and Pattern Recognition*, pp. 770-778, 2016.

[7] A. Vaswani et al., "Attention Is All You Need," *Advances in Neural Information Processing Systems*, pp. 5998-6008, 2017.

[8] X. Shi, Z. Chen, H. Wang, D.-Y. Yeung, W.-K. Wong, and W.-C. Woo, "Convolutional LSTM Network: A Machine Learning Approach for Precipitation Nowcasting," *Advances in Neural Information Processing Systems*, pp. 802-810, 2015.

[9] Y. Wang, M. Long, J. Wang, Z. Gao, and P. S. Yu, "PredRNN: Recurrent Neural Networks for Predictive Learning using Spatiotemporal LSTMs," *Advances in Neural Information Processing Systems*, pp. 879-888, 2017.

[10] J. Zhang, Y. Zheng, and D. Qi, "Deep Spatio-Temporal Residual Networks for Citywide Crowd Flows Prediction," *AAAI Conference on Artificial Intelligence*, pp. 1655-1661, 2018.

[11] D. P. Kingma and J. Ba, "Adam: A Method for Stochastic Optimization," *International Conference on Learning Representations*, 2015.

[12] Y. Gal and Z. Ghahramani, "Dropout as a Bayesian Approximation: Representing Model Uncertainty in Deep Learning," *International Conference on Machine Learning*, pp. 1050-1059, 2016.

[13] M. Pesaresi et al., "GHS-BUILT-S R2023A - GHS Built-Up Surface Grid, derived from Sentinel2 composite and Landsat, multitemporal (1975-2030)," *European Commission, Joint Research Centre (JRC)*, 2023. [Online]. Available: http://data.europa.eu/89h/9f06f36f-4b11-47ec-abb0-4f8b7b1d72ea